

# Efficient Reinforcement Learning with Hierarchies of Machines by Leveraging Internal Transitions

Aijun Bai\*

UC Berkeley/Microsoft Research

Stuart Russell


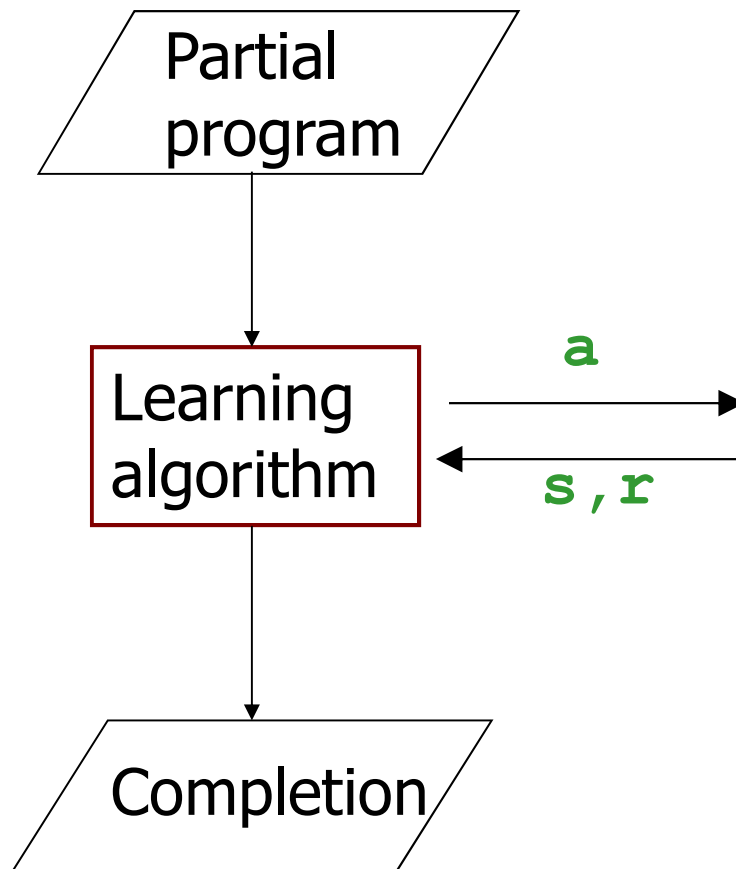
UC Berkeley

# Outline

- Hierarchical RL with partial programs
- Deterministic internal transitions
- Results

# Hierarchical RL with partial programs

[Parr & Russell, NIPS 97; Andre & Russell, NIPS 00, AAAI 02; Marthi et al, IJCAI 05]



*Hierarchically optimal*  
for all terminating programs

# Partial Program – an Example

repeat forever

Choose({a1,a2,...})

# Partial Program – an Example

Navigate(destination)

while  $\neg \text{At}(\text{destination}, \text{CurrentState}())$

Choose({N,S,E,W})

# Concurrent Partial Programs

Top()

  for each p in **Effectors**()

    PlayKeep(p)

PlayKeep(p)

  s  $\leftarrow$  CurrentState()

  while  $\neg$ Terminal(s)

    if BallKickable(s) then **Choose**({Pass(),Hold()})

    else if FastestToBall(s) then Intercept()

    else **Choose**(Stay(),Move())

Pass()

  KickTo(**Choose**(Effectors()\{self}),**Choose**({slow,fast}))

...

# Technical development

- Decisions based on *internal state*
  - Joint state  $\omega = [s, m]$  environment state + program state (cf. [Russell & Wefald 1989] )
- MDP + partial program = SMDP over choice states in  $\{\omega\}$ , learn  $Q(\omega, c)$  for choices  $c$
- Additive *decomposition* of value functions
  - by *subroutine structure* [Dietterich 00, Andre & Russell 02]  
 $Q$  is a sum of sub- $Q$  functions per subroutine
  - across *concurrent threads* [Russell & Zimdars 03]  
 $Q$  is a sum of sub- $Q$  functions per thread, with decomposed reward signal

# Internal Transitions

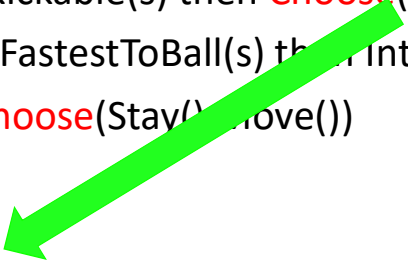
- Transitions between choice points with no physical action intervening
- Internal transitions take no (real) time and have zero reward
- Internal transitions are deterministic

```
Top()
  for each p in Effectors()
    PlayKeep(p)

PlayKeep(p)
  s ← CurrentState()
  while ¬Terminal(s)
    if BallKickable(s) then Choose({Pass(),Hold()})
    else if FastestToBall(s) then Intercept()
    else Choose(Stay(),Move())

Pass()
  KickTo(Choose(Effectors()\{self}),Choose({slow,fast}))

...
```





# Idea 1

- Use internal transitions to shortcircuit the computations of Q values recursively if applicable
  - If  $(s, m, c) \rightarrow (s, m')$  is an internal transition
  - Then,  $Q(s, m, c) = V(s, m') = \max_{c'} Q(s, m', c')$
- Cache internal transitions as **<s, m, c, m'>** tuples
- No need for Q-learning on these

## Idea 2

- Identify weakest precondition  $P(s)$  for this internal transition to occur (cf EBL, chunking)
- Cache internal transitions as  $\langle \mathbf{P}, \mathbf{m}, \mathbf{c}, \mathbf{m}' \rangle$  tuples
- Cache size independent of  $|S|$ , roughly proportional to size of partial program call graph

# The HAMQ-INT Algorithm

- Track the set of predicates since last choice point
- Save an abstracted rule of internal transition if qualified ( $\tau = 0$ ) in a dictionary  $\rho$
- Use the saved rules to shortcircuit the computation of  $Q$  values recursively whenever possible


```

QUpdate ( $s' : state, z' : stack, r : reward,$ 
           $t' : current\ time, \mathcal{P} : evaluated\ predicates$ ) :
  if  $t' = t$  then
     $\rho[\mathcal{P}, \mathcal{P}(s), z, c] \leftarrow z'$ 
  else
     $\mathbf{QTable}(s, z, c) \leftarrow (1 - \alpha) \mathbf{QTable}(s, z, c)$ 
     $\quad + \alpha(r + \gamma^{t'-t} \max_{c'} \mathbf{Q}(s', z', c'))$ 
     $(t, s, z) \leftarrow (t', s', z')$ 

 $\mathbf{Q}(s : state, z : stack, c : choice)$  :
  if  $\exists \mathcal{P} \text{ s.t. } \langle \mathcal{P}, \mathcal{P}(s), z, c \rangle \in \rho.\mathbf{Keys}()$  then
     $q \leftarrow -\infty$ 
     $z' \leftarrow \rho[\mathcal{P}, \mathcal{P}(s), z, c]$ 
    for  $c' \in \mu(z.\mathbf{Top}())$  do
       $q \leftarrow \max(q, \mathbf{Q}(s, z', c'))$ 
    return  $q$ 
  else
    return  $\mathbf{QTable}(s, z, c)$ 

```


# Experimental Result on Taxi




# 3 vs 2 Keepaway Comparisons

- Option (Stone, 2005):
  - Each keeper learning separately
  - Learn a policy over Hold() and Pass(k, v) if ball kickable; otherwise, follow a fixed policy
    - Intercept() if fastest to the ball; otherwise, GetOpen()
    - GetOpen() is manually programmed for Option
- Concurrent-Option:
  - Concurrent version of Option
    - One global Q function is learnt
- Random: randomized version of Option
- Concurrent-HAMQ
  - Learn its own version of GetOpen() by calling Stay() and Move(d, v)
- Concurrent-HAMQ-INT


# Experimental Result on Keepaway



# Before and After



Initial policy



Converged policy

# Summary

- HAMQ-INT algorithm
  - Automatically discovers internal transitions
  - Takes advantage of internal transitions for efficient learning
  - Outperforms the state of the art significantly on Taxi and RoboCup Keepaway
- Future work
  - Scale up to full RoboCup task
  - More general integration of model-based and model-free reinforcement learning
  - More flexible forms of partial program (e.g., temporal logic)